





v1: Learning to Point Visual Tokens for Multimodal Grounded Reasoning

First authors
Jiwan Chung^{1,*} Junhyeok Kim^{1,*} Siyeol Kim¹ Jaeyoung Lee¹ Min Soo Kim¹ Youngjae Yu²

¹ Yonsei University ² Seoul National University · * Equal contribution

MOTIVATION

**Visual attention fades during reasoning.
v1 explicitly revisits relevant patches.**



Find the value of z .


Text-based reasoning can attend to image tokens, but visual attention fades

"... In an isosceles right triangle, the two equal sides would be $2z - 15 = 7 \rightarrow z = 11$." ✗

v1 points to and copies image regions throughout reasoning


"...detect(query="angle at S") <region>  - 15 </region> the angles at S and T seem to be the same, so $RS = RT$. Thus $2z - 15 = 9 \rightarrow z = 12$." ✓

TVC-7B on RefCOCO testA · attention from the newest token to the image tokens · layers 2, 14, 27 · paper Fig. 3



Total attention on image tokens per step.

Declines steadily throughout decoding.




Relative target-region attention: target-box mean / all-image-token mean.

Converges to ≈ 0.8 by mid-generation, below average.

METHOD • POINT-AND-COPY


Generate a word, or point to a patch and copy it.



$$\text{logit}_{\text{ptr}}^{(k)} = \mathbf{L}_q(\mathbf{h}_t) \cdot \mathbf{L}_k(\mathbf{c}_k)^\top / \sqrt{D}$$

- **Pointing logit** $\text{logit}_{\text{ptr}}^{(k)}$: copy-head score for patch k , a query from the hidden state h_l against a key from the patch embedding c_k ; L_q, L_k are the only new parameters.
- **Training** • **v1q**. 300K reasoning traces with interleaved grounding annotations, on Qwen2.5-VL-7B.

QUALITATIVE EXAMPLE



Let's see here. I've got this bar graph.

```
detect(query="bar graph titled \"Sweet Tooth\"", objects=["<obj1><region>[0.23, 0.25, 0.96, 0.79]</copy_76> ... </copy_287></region>"])
```

The bar graph **<obj1>** shows the number of students labeled from A to E.

... (y-axis look-up omitted) ...

```
detect(query="bar for Candy E", objects=["<obj8><region>[0.74, 0.37, 0.96, 0.88]
```

```
<|copy 139|> ... <|copy 303|></region>" )
```

Candy E <obj8> has a bar reaching up to 5.
There are 25 students in total. So, percentage = $(5 / 25) * 100 = 20\%$

RESULTS · MULTIMODAL MATH REASONING

Outperforms 7B baselines; narrows the gap to 72B models.

Model	Size	Reasoning Only	MathVista mini	MathVision full	MathVerse mini	Average mini	Average full
Qwen2-VL (Wang et al., 2024b)	7B	✗	60.9	-	16.3	24.6	-
Qwen2-VL (Wang et al., 2024b)	72B	✗	69.7	-	26.6	36.2	-
Qwen2.5-VL (Bai et al., 2025)	7B	✗	67.8	23.6	-	44.5	45.3
Qwen2.5-VL (Bai et al., 2025)	72B	✗	74.8	39.8	-	57.6	57.4
InternVL2.5 (Chen et al., 2025)	8B	✗	64.4	22.0	19.7	39.5	41.9
InternVL2.5 (Chen et al., 2025)	78B	✗	72.3	34.9	32.2	51.7	53.0
GPT-4o (Hurst et al., 2024)	-	✗	63.8	-	30.4	50.2	-
LLaVa-CoT (Xu et al., 2025)	11B	✓	54.8	16.3	-	33.9	35.0
Mulberry (Yao et al., 2024)	7B	✓	63.1	-	-	39.6	-
TVC (Sun et al., 2025)	7B	✓	68.1	-	22.7	38.9	-
TVC (Sun et al., 2025)	72B	✓	72.2	-	41.9	48.8	-
QVQ-72B-preview (Qwen Team, 2024)	72B	✓	71.4	35.9	-	41.5	49.6
Base (Qwen2.5-VL)	7B	✗	67.8	23.6	-	44.5	45.3
Text-Only (TVC)	7B	✓	68.1	-	22.7	38.9	-
Ours	7B	✓	68.6	34.5	28.1	48.6	50.6
↳ Inference w/o Pointing	7B	✓	60.0	25.3	23.7	33.6	28.9

Table 1 · accuracy (%), Gemini-2.0-Flash judge · bold: best 7B-scale result · 72B/78B and GPT-4o are reference points

ANALYSIS

Variant	Tr.	Inf.	Score	Δ
Backbone	✗	✗	23.6	–
(no pointing)	✓	✗	25.3	+1.
Ours	✓	✓	34.5	+10.

Pointing at inference drives the gain.

Disabling point-and-copy at inference drops MathVision-mini accuracy from 34.5 to 25.3.

v1 sees the image again: attention lands on the copied patches.




Fig. 5 · attention to copied vs. original visual tokens after the first copy. Layers 2 and 14 attend mostly to the copied tokens; layer 27 balances copied and original tokens.